

Summary of EDA of the Ames Housing Dataset

Background

This week, we explored Kaggle's "House Prices: Advanced Regression Techniques" competition to sharpen our Exploratory Data Analysis (EDA) skills. Our focus is on house prices in Ames, Iowa, represented by the dependent variable SalePrice.

Management/Research Question

The research question is: "What factors influence house prices in Ames, Iowa?" Understanding these factors can help buyers, sellers, and real estate professionals make informed decisions.

Requirements and Methods

We conducted EDA on the dataset as follows:

1. **Descriptive Statistics and Visualizations:** Analyzed the marginal distribution of SalePrice.
2. **Missing Data and Outliers:** Investigated missing values and outliers.
3. **Potential Predictors:** Examined three potential predictors of SalePrice with graphs and statistics.
4. **Feature Creation:** Generated a new predictor through feature engineering.
5. **Scaling Methods:** Applied min-max and standard scaling to the dependent variable.

Descriptive Statistics and Visualizations

The dataset contains 1,460 observations and 81 variables. SalePrice ranges from \$34,900 to \$755,000, with a mean of \$180,921. The distribution is right-skewed, suggesting a log transformation for normalization.

Missing Data and Outliers

Nineteen variables have missing values, with Alley, PoolQC, Fence, and MiscFeature having over 50% missingness and subsequently removed. Boxplots identified several outliers, but these values were retained as they were realistic.

Potential Predictors

We investigated the relationships between SalePrice and three predictors:

1. **GrLivArea:** Shows a strong positive correlation with SalePrice, indicating that larger living areas generally command higher prices.
2. **OverallQual:** Another strong predictor; higher quality ratings correlate with higher prices.
3. **GarageCars:** Homes with more garage spaces tend to have higher sale prices.

The correlations are as follows:

- GrLivArea: $r = 0.71$ $r = 0.71$
- OverallQual: $r = 0.79$ $r = 0.79$
- GarageCars: $r = 0.62$ $r = 0.62$

Feature Creation

We created the feature TotalSF, which is the sum of GrLivArea and TotalBsmtSF. This new feature showed a strong association with SalePrice ($r = 0.81$ $r = 0.81$).

Scaling Methods

Both min-max and standard scaling were applied to SalePrice:

- **Min-Max Scaling:** Adjusted SalePrice to a 0-1 range.
- **Standard Scaling:** Centered SalePrice around the mean with a standard deviation of 1.

Results and Insights

1. **GrLivArea** and SalePrice have a strong correlation, but two larger homes deviate, indicating potential outliers.
2. **OverallQual** shows a clear trend where higher quality ratings significantly increase the median SalePrice, from \$50,150 for quality 1 to \$432,390 for quality 10.
3. **GarageCars:** As the number of garage spaces increases, both median and mean SalePrice rise, with homes having 3 garage spaces showing the highest median SalePrice of \$295,000.

Conclusion

Our EDA provides valuable insights into the factors influencing house prices in Ames, Iowa. The relationship between GrLivArea, OverallQual, and GarageCars with SalePrice highlights the importance of these features. Feature engineering and scaling techniques further refined our understanding, paving the way for more advanced modeling.

In summary, this EDA revealed insights about factors that may influence the home prices in Ames, Iowa. Our parsimonious regression model that used three predictor variables (TotalSF, YrSinceRemod, GarageCar) accounted for roughly 75% of the variation in Sale Price. Further research and a more sophisticated framework is needed to explain and elucidate the remaining 25% of variation. This baseline regression model can benefit homebuyers and sellers aiming to better understand fair market prices for homes in Ames, Iowa. The lack of ability to study the distribution of the error terms is a limitation of our approach. Other avenues to pursue would include using principal components or clustering techniques on the 80 variables to leverage the dimensionality of the dataset.

Appendix:

Module 1 Assignment 1: House Prices: Advanced Regression Techniques EDA (Kaggle)

Sachin Sharma

MSDS 422

2024/06/23

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
import scipy.stats as st
import seaborn as sns

import warnings

# Ignore all FutureWarnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

Load Train Data

```
train_df = pd.read_csv("train.csv")
train_df.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
5	Lvl	AllPub	...	0	NaN	NaN	NaN	0
9	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0

```

2
4      Lvl    AllPub    ...      0    NaN    NaN      NaN      0
12

   YrSold  SaleType  SaleCondition  SalePrice
0    2008         WD         Normal    208500
1    2007         WD         Normal    181500
2    2008         WD         Normal    223500
3    2006         WD        Abnorml    140000
4    2008         WD         Normal    250000

[5 rows x 81 columns]

print('Columns: \n'+", ".join(train_df.columns.tolist()))

Columns:
Id, MSSubClass, MSZoning, LotFrontage, LotArea, Street, Alley,
LotShape, LandContour, Utilities, LotConfig, LandSlope, Neighborhood,
Condition1, Condition2, BldgType, HouseStyle, OverallQual,
OverallCond, YearBuilt, YearRemodAdd, RoofStyle, RoofMatl,
Exterior1st, Exterior2nd, MasVnrType, MasVnrArea, ExterQual,
ExterCond, Foundation, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1,
BsmtFinSF1, BsmtFinType2, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, Heating,
HeatingQC, CentralAir, Electrical, 1stFlrSF, 2ndFlrSF, LowQualFinSF,
GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath,
BedroomAbvGr, KitchenAbvGr, KitchenQual, TotRmsAbvGrd, Functional,
Fireplaces, FireplaceQu, GarageType, GarageYrBlt, GarageFinish,
GarageCars, GarageArea, GarageQual, GarageCond, PavedDrive,
WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch,
PoolArea, PoolQC, Fence, MiscFeature, MiscVal, MoSold, YrSold,
SaleType, SaleCondition, SalePrice

train_df.shape

(1460, 81)

```

EDA

1. Provide appropriate descriptive statistics and visualizations to help understand the marginal distribution of the dependent variable.

The dependent variable is SalePrice. There are 1460 entries in the training data

```

train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459

```

```

Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     1460 non-null   int64
1   MSSubClass             1460 non-null   int64
2   MSZoning               1460 non-null   object
3   LotFrontage            1201 non-null   float64
4   LotArea                1460 non-null   int64
5   Street                 1460 non-null   object
6   Alley                  91 non-null     object
7   LotShape               1460 non-null   object
8   LandContour            1460 non-null   object
9   Utilities              1460 non-null   object
10  LotConfig              1460 non-null   object
11  LandSlope              1460 non-null   object
12  Neighborhood           1460 non-null   object
13  Condition1             1460 non-null   object
14  Condition2             1460 non-null   object
15  BldgType               1460 non-null   object
16  HouseStyle             1460 non-null   object
17  OverallQual            1460 non-null   int64
18  OverallCond            1460 non-null   int64
19  YearBuilt              1460 non-null   int64
20  YearRemodAdd           1460 non-null   int64
21  RoofStyle              1460 non-null   object
22  RoofMatl               1460 non-null   object
23  Exterior1st            1460 non-null   object
24  Exterior2nd            1460 non-null   object
25  MasVnrType             588 non-null    object
26  MasVnrArea             1452 non-null   float64
27  ExterQual              1460 non-null   object
28  ExterCond              1460 non-null   object
29  Foundation             1460 non-null   object
30  BsmtQual               1423 non-null   object
31  BsmtCond               1423 non-null   object
32  BsmtExposure           1422 non-null   object
33  BsmtFinType1           1423 non-null   object
34  BsmtFinSF1            1460 non-null   int64
35  BsmtFinType2           1422 non-null   object
36  BsmtFinSF2            1460 non-null   int64
37  BsmtUnfSF             1460 non-null   int64
38  TotalBsmtSF            1460 non-null   int64
39  Heating                1460 non-null   object
40  HeatingQC              1460 non-null   object
41  CentralAir             1460 non-null   object
42  Electrical             1459 non-null   object
43  1stFlrSF               1460 non-null   int64
44  2ndFlrSF               1460 non-null   int64
45  LowQualFinSF           1460 non-null   int64

```

```

46 GrLivArea      1460 non-null int64
47 BsmtFullBath   1460 non-null int64
48 BsmtHalfBath   1460 non-null int64
49 FullBath       1460 non-null int64
50 HalfBath       1460 non-null int64
51 BedroomAbvGr  1460 non-null int64
52 KitchenAbvGr  1460 non-null int64
53 KitchenQual    1460 non-null object
54 TotRmsAbvGrd  1460 non-null int64
55 Functional     1460 non-null object
56 Fireplaces     1460 non-null int64
57 FireplaceQu    770 non-null object
58 GarageType     1379 non-null object
59 GarageYrBlt    1379 non-null float64
60 GarageFinish   1379 non-null object
61 GarageCars     1460 non-null int64
62 GarageArea     1460 non-null int64
63 GarageQual     1379 non-null object
64 GarageCond     1379 non-null object
65 PavedDrive     1460 non-null object
66 WoodDeckSF     1460 non-null int64
67 OpenPorchSF    1460 non-null int64
68 EnclosedPorch  1460 non-null int64
69 3SsnPorch      1460 non-null int64
70 ScreenPorch    1460 non-null int64
71 PoolArea       1460 non-null int64
72 PoolQC         7 non-null object
73 Fence          281 non-null object
74 MiscFeature    54 non-null object
75 MiscVal        1460 non-null int64
76 MoSold         1460 non-null int64
77 YrSold         1460 non-null int64
78 SaleType       1460 non-null object
79 SaleCondition  1460 non-null object
80 SalePrice      1460 non-null int64

```

dtypes: float64(3), int64(35), object(43)

memory usage: 924.0+ KB

train_df.describe()

	Id	MSSubClass	LotFrontage	LotArea
OverallQual \				
count	1460.000000	1460.000000	1201.000000	1460.000000
1460.000000				
mean	730.500000	56.897260	70.049958	10516.828082
6.099315				
std	421.610009	42.300571	24.284752	9981.264932
1.382997				
min	1.000000	20.000000	21.000000	1300.000000
1.000000				

25%	365.750000	20.000000	59.000000	7553.500000
5.000000				
50%	730.500000	50.000000	69.000000	9478.500000
6.000000				
75%	1095.250000	70.000000	80.000000	11601.500000
7.000000				
max	1460.000000	190.000000	313.000000	215245.000000
10.000000				

	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
BsmtFinSF1	...			
count	1460.000000	1460.000000	1460.000000	1452.000000
1460.000000	...			
mean	5.575342	1971.267808	1984.865753	103.685262
443.639726	...			
std	1.112799	30.202904	20.645407	181.066207
456.098091	...			
min	1.000000	1872.000000	1950.000000	0.000000
0.000000	...			
25%	5.000000	1954.000000	1967.000000	0.000000
0.000000	...			
50%	5.000000	1973.000000	1994.000000	0.000000
383.500000	...			
75%	6.000000	2000.000000	2004.000000	166.000000
712.250000	...			
max	9.000000	2010.000000	2010.000000	1600.000000
5644.000000	...			

	WoodDeckSF	OpenPorchSF	EnclosedPorch	3SsnPorch
ScreenPorch				
count	1460.000000	1460.000000	1460.000000	1460.000000
1460.000000				
mean	94.244521	46.660274	21.954110	3.409589
15.060959				
std	125.338794	66.256028	61.119149	29.317331
55.757415				
min	0.000000	0.000000	0.000000	0.000000
0.000000				
25%	0.000000	0.000000	0.000000	0.000000
0.000000				
50%	0.000000	25.000000	0.000000	0.000000
0.000000				
75%	168.000000	68.000000	0.000000	0.000000
0.000000				
max	857.000000	547.000000	552.000000	508.000000
480.000000				

	PoolArea	MiscVal	MoSold	YrSold
SalePrice				
count	1460.000000	1460.000000	1460.000000	1460.000000

```

1460.000000
mean      2.758904      43.489041      6.321918      2007.815753
180921.195890
std       40.177307     496.123024      2.703626      1.328095
79442.502883
min        0.000000      0.000000      1.000000      2006.000000
34900.000000
25%        0.000000      0.000000      5.000000      2007.000000
129975.000000
50%        0.000000      0.000000      6.000000      2008.000000
163000.000000
75%        0.000000      0.000000      8.000000      2009.000000
214000.000000
max       738.000000    15500.000000     12.000000     2010.000000
755000.000000

[8 rows x 38 columns]

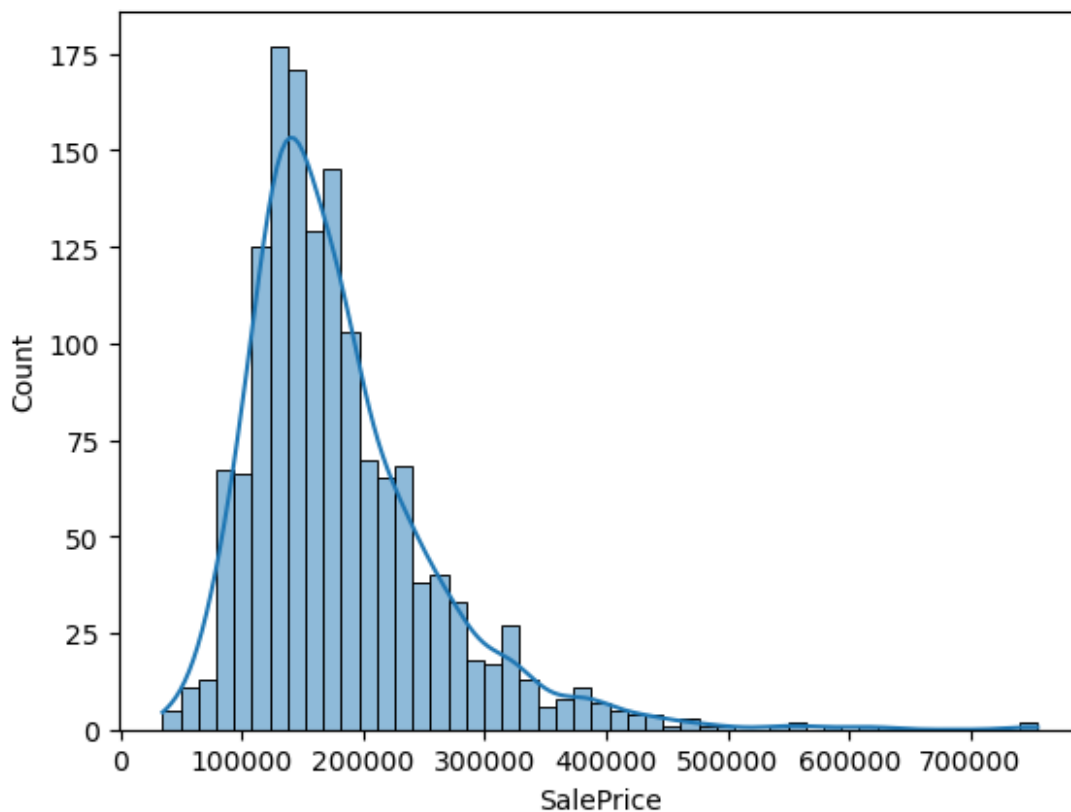
```

Distribution of SalePrice:

```

sns.histplot(data=train_df, kde=True, x='SalePrice')
<Axes: xlabel='SalePrice', ylabel='Count'>

```

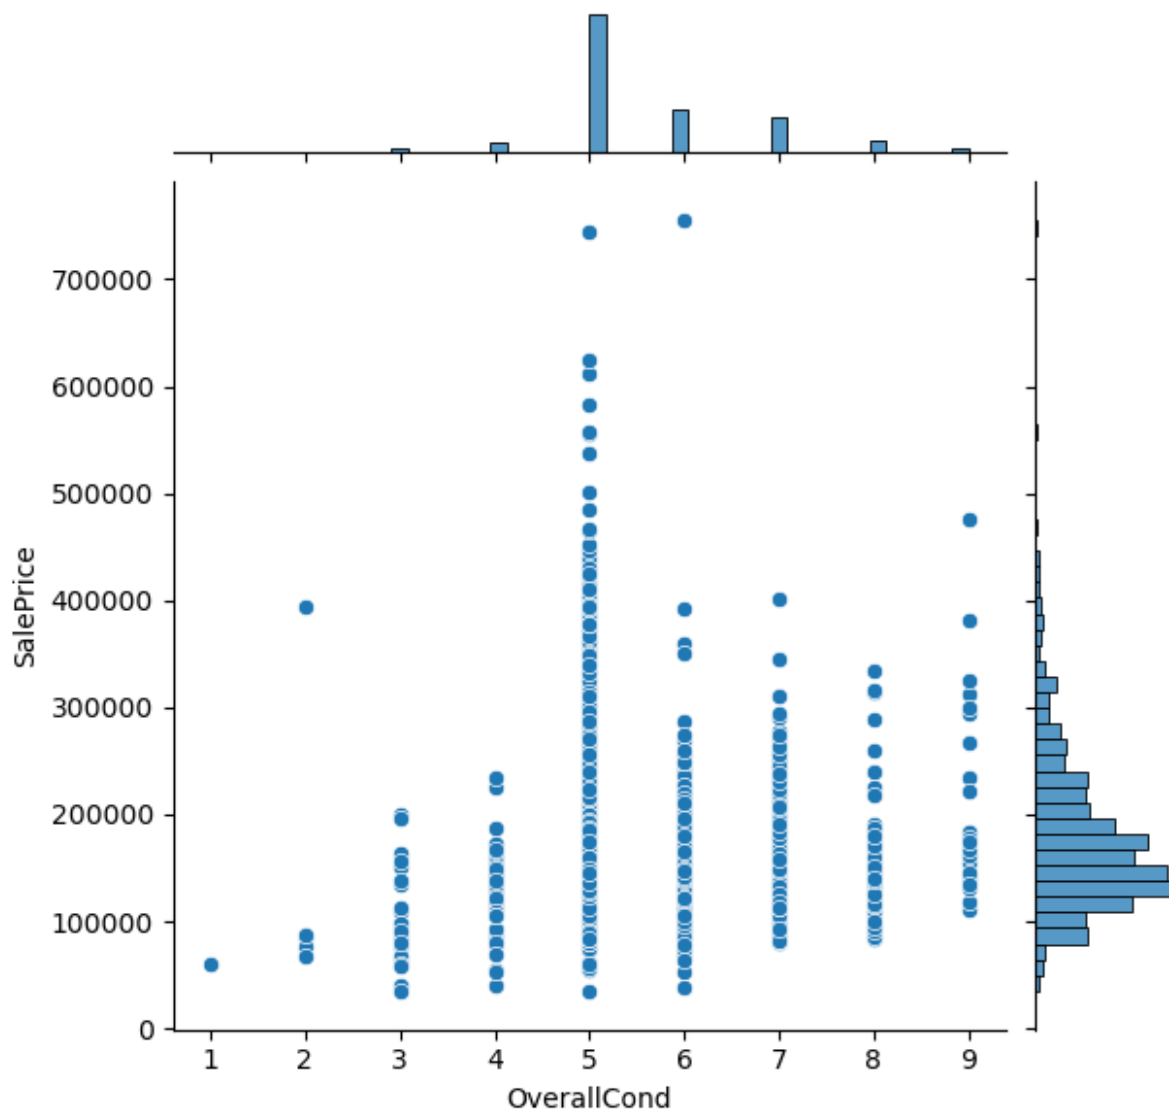


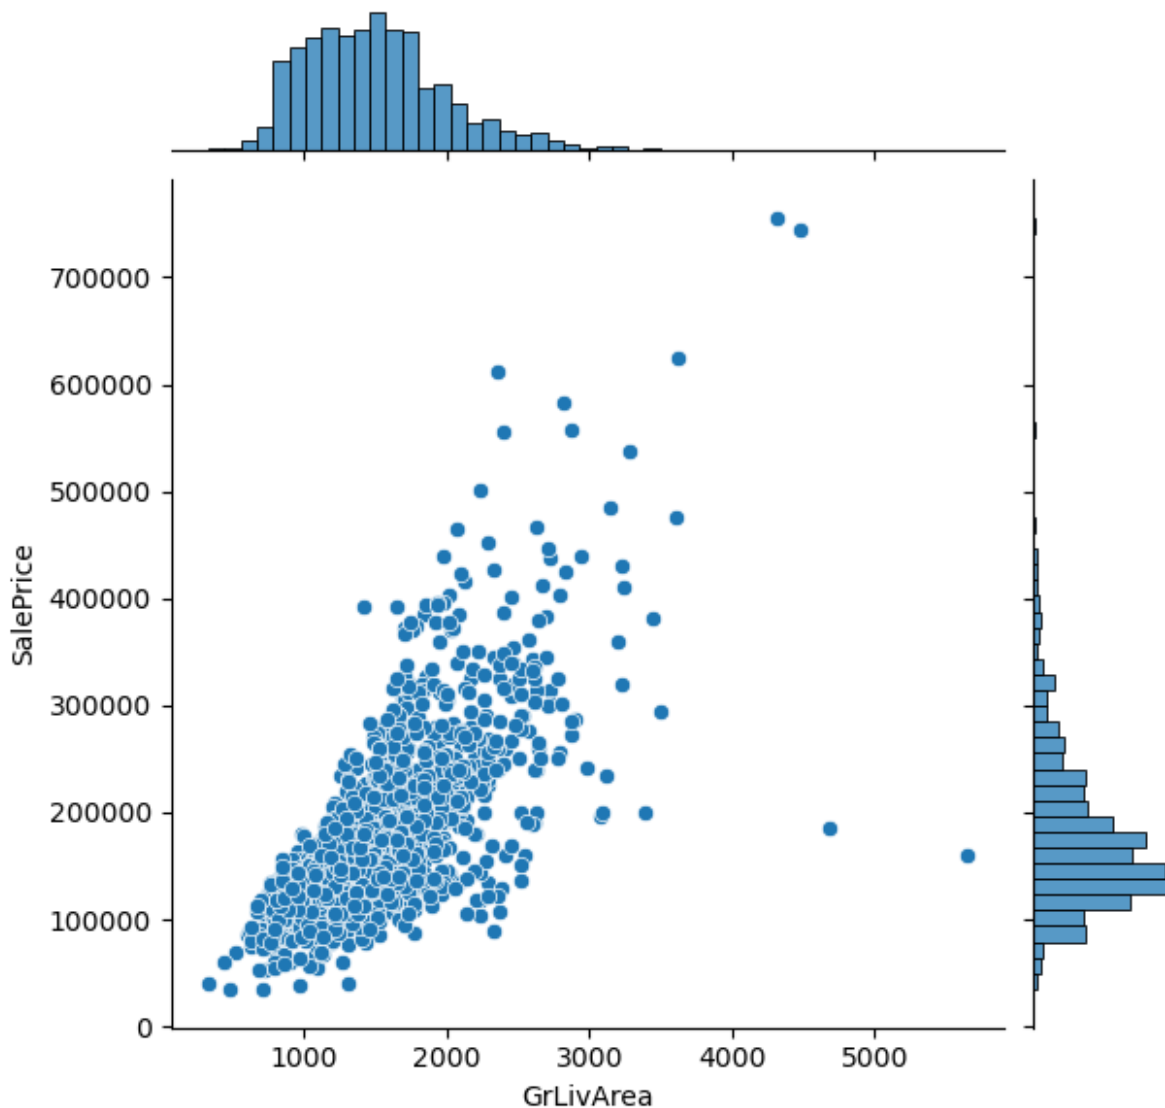

```
train_df.SalePrice.describe()
count      1460.000000
mean       180921.195890
std         79442.502883
min         34900.000000
25%        129975.000000
50%        163000.000000
75%        214000.000000
max         755000.000000
Name: SalePrice, dtype: float64
```

The median SalePrice is \ \$163000 while the mean is ~\ \$180921.20.

Joint Plots of SalePrice in respect to other features like OverallCond, GrLivingArea:

```
sns.jointplot(x="OverallCond",
y="SalePrice",edgecolor="white",data=train_df);
sns.jointplot(x="GrLivArea",
y="SalePrice",edgecolor="white",data=train_df);
```





Qualitative/categorical and Quantitative/numerical data:

```
# Get a list of qualitative columns in the training set, where column
type is 'object'
qual_cols = train_df.select_dtypes(include='object').columns.tolist()
print('Qualitative Columns: \n' + ", ".join(qual_cols))

# Get a list of quantitative columns in the training set, where column
type is not 'object'
quant_cols = train_df.select_dtypes(exclude='object').columns.tolist()
print('\nQuantitative Columns: \n' + ", ".join(quant_cols))

Qualitative Columns:
MSZoning, Street, Alley, LotShape, LandContour, Utilities, LotConfig,
LandSlope, Neighborhood, Condition1, Condition2, BldgType, HouseStyle,
RoofStyle, RoofMatl, Exterior1st, Exterior2nd, MasVnrType, ExterQual,
```

ExterCond, Foundation, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, Heating, HeatingQC, CentralAir, Electrical, KitchenQual, Functional, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, PavedDrive, PoolQC, Fence, MiscFeature, SaleType, SaleCondition

Quantitative Columns:

Id, MSSubClass, LotFrontage, LotArea, OverallQual, OverallCond, YearBuilt, YearRemodAdd, MasVnrArea, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, 1stFlrSF, 2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, TotRmsAbvGrd, Fireplaces, GarageYrBlt, GarageCars, GarageArea, WoodDeckSF, OpenPorchSF, EnclosedPorch, 3SsnPorch, ScreenPorch, PoolArea, MiscVal, MoSold, YrSold, SalePrice

2. Investigate missing data and outliers.

```
nullTotals = train_df.isnull().sum().sort_values(ascending = False)
percentageOfNull = (train_df.isnull().sum() /
train_df.isnull().count()).sort_values(ascending = False)
emptyVals = pd.concat([nullTotals, percentageOfNull], axis=1,
keys=['Total Missing Values', 'Percentage Null'])
emptyVals.head(20)
```

	Total Missing Values	Percentage Null
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
MasVnrType	872	0.597260
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageYrBlt	81	0.055479
GarageCond	81	0.055479
GarageType	81	0.055479
GarageFinish	81	0.055479
GarageQual	81	0.055479
BsmtFinType2	38	0.026027
BsmtExposure	38	0.026027
BsmtQual	37	0.025342
BsmtCond	37	0.025342
BsmtFinType1	37	0.025342
MasVnrArea	8	0.005479
Electrical	1	0.000685
Id	0	0.000000

```
num_na_features = train_df.isna().sum()
num_na_features = num_na_features[num_na_features > 0]
num_na_features.sort_values(inplace=True)
num_na_features.describe()
```

```
count      19.000000
mean       412.052632
std        551.246917
min         1.000000
25%        37.500000
50%        81.000000
75%       781.000000
max      1453.000000
dtype: float64
```

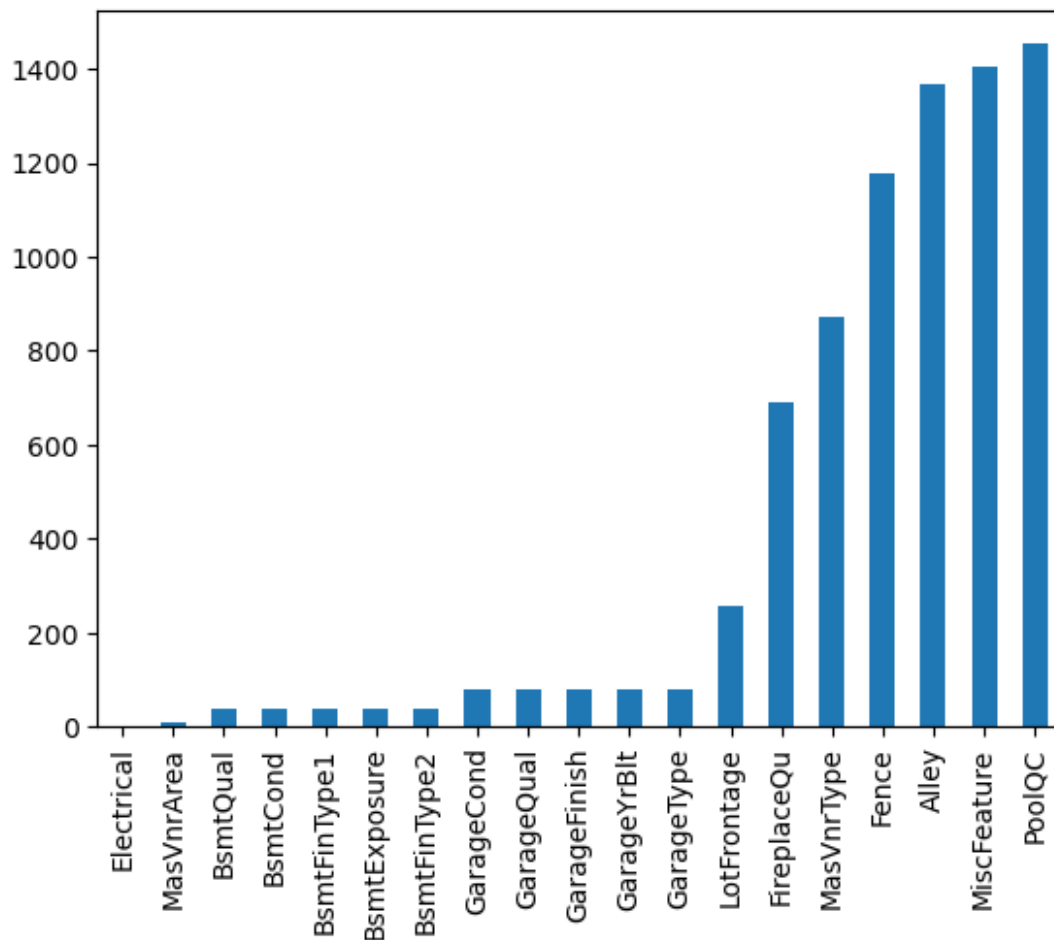
```
num_na_features.size
```

```
19
```

Let's graph a bar chart of null values.

```
num_na_features.plot.bar()
```

```
<Axes: >
```



There are 19 features that contain null values.

Most listings/rows do not contain 'PoolQC', 'MiscFeature', 'Alley', 'Fence', 'FireplaceQC', etc features.

Almost half of the homes do not list 'FireplaceQC' (690)

```
train_df[train_df['Fireplaces'] == 0].FireplaceQu.isnull().sum()  
690
```

All but 7 of the 1453 listings/rows do not have 'PoolQc'

```
len(train_df[train_df['PoolArea'] != 0])  
7
```

1369 null in Alley column

```
train_df['Alley'].isnull().sum()  
1369
```

Let's examine the category outliers with respect to the dependent variable, SalePrice. The total number of outliers within SalePrice is 61.

```
# get a list of qualitative columns in the training set, where column  
type = 'object'  
qual_cols = [f for f in train_df.columns if train_df.dtypes[f] ==  
'object']  
qual_cols.append('SalePrice') # add the quantitative dependent  
variable (SalePrice) to the list  
train_qual_df = train_df[qual_cols] # get a dataframe of qualitative  
columns mapped to dependent variable (SalePrice)  
train_qual_df = train_qual_df.fillna('NONE') # fill any null cells in  
the dataframe with NONE  
  
Q1 = train_df['SalePrice'].quantile(0.25)  
Q3 = train_df['SalePrice'].quantile(0.75)  
IQR = Q3 - Q1  
totalOutliers = ((train_df['SalePrice'] < (Q1 - 1.5 * IQR)) |  
(train_df['SalePrice'] > (Q3 + 1.5 * IQR))).sum()  
print("IQR value: {}  
Total amount of outliers within SalePrice:  
{}".format(IQR, totalOutliers))  
  
IQR value: 84025.0  
Total amount of outliers within SalePrice: 61
```

3. Investigate at least three potential predictors of the dependent variable and provide appropriate graphs / statistics to demonstrate the relationships.

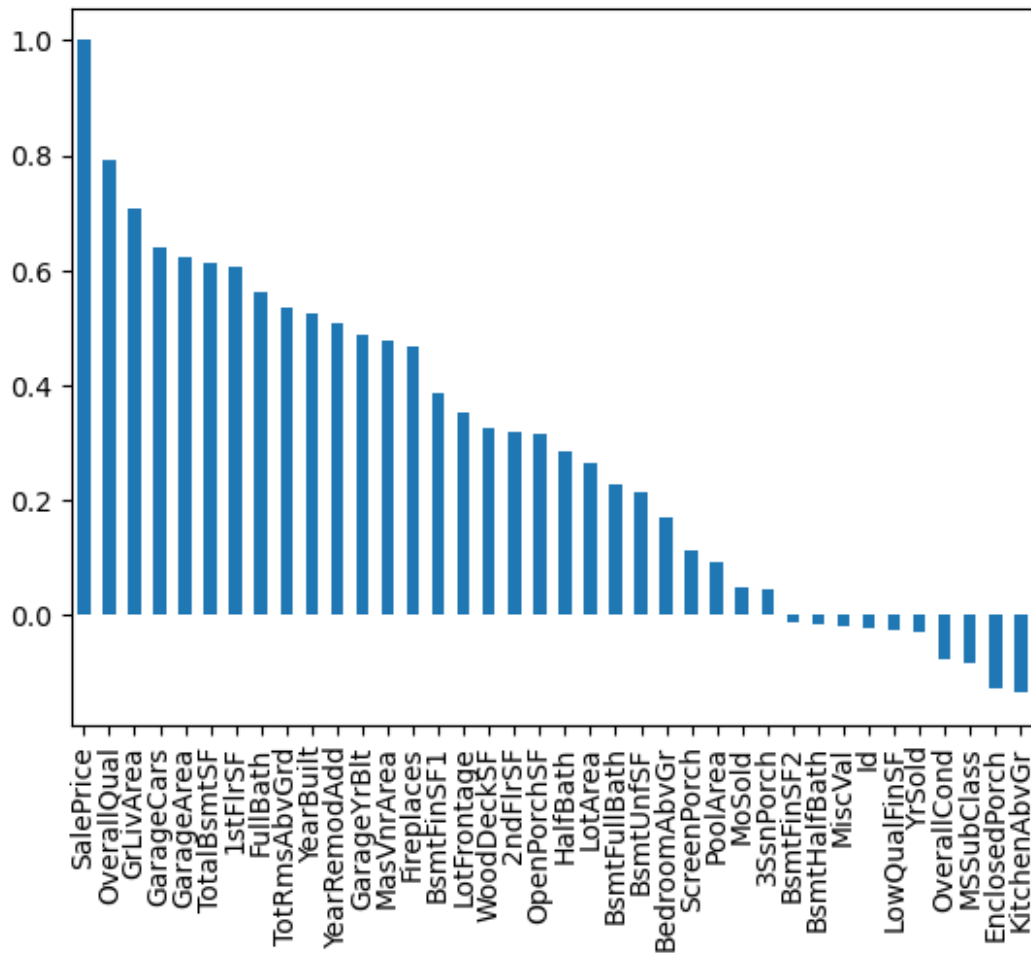
```
# Calculate correlation matrix for numeric columns only
corr_matrix = train_df.corr(method='pearson', min_periods=30,
numeric_only=True)
corr_matrix_sorted =
corr_matrix['SalePrice'].sort_values(ascending=False)
print(corr_matrix_sorted)
```

SalePrice	1.000000
OverallQual	0.790982
GrLivArea	0.708624
GarageCars	0.640409
GarageArea	0.623431
TotalBsmtSF	0.613581
1stFlrSF	0.605852
FullBath	0.560664
TotRmsAbvGrd	0.533723
YearBuilt	0.522897
YearRemodAdd	0.507101
GarageYrBlt	0.486362
MasVnrArea	0.477493
Fireplaces	0.466929
BsmtFinSF1	0.386420
LotFrontage	0.351799
WoodDeckSF	0.324413
2ndFlrSF	0.319334
OpenPorchSF	0.315856
HalfBath	0.284108
LotArea	0.263843
BsmtFullBath	0.227122
BsmtUnfSF	0.214479
BedroomAbvGr	0.168213
ScreenPorch	0.111447
PoolArea	0.092404
MoSold	0.046432
3SsnPorch	0.044584
BsmtFinSF2	-0.011378
BsmtHalfBath	-0.016844
MiscVal	-0.021190
Id	-0.021917
LowQualFinSF	-0.025606
YrSold	-0.028923
OverallCond	-0.077856
MSSubClass	-0.084284
EnclosedPorch	-0.128578

```
KitchenAbvGr      -0.135907  
Name: SalePrice, dtype: float64
```

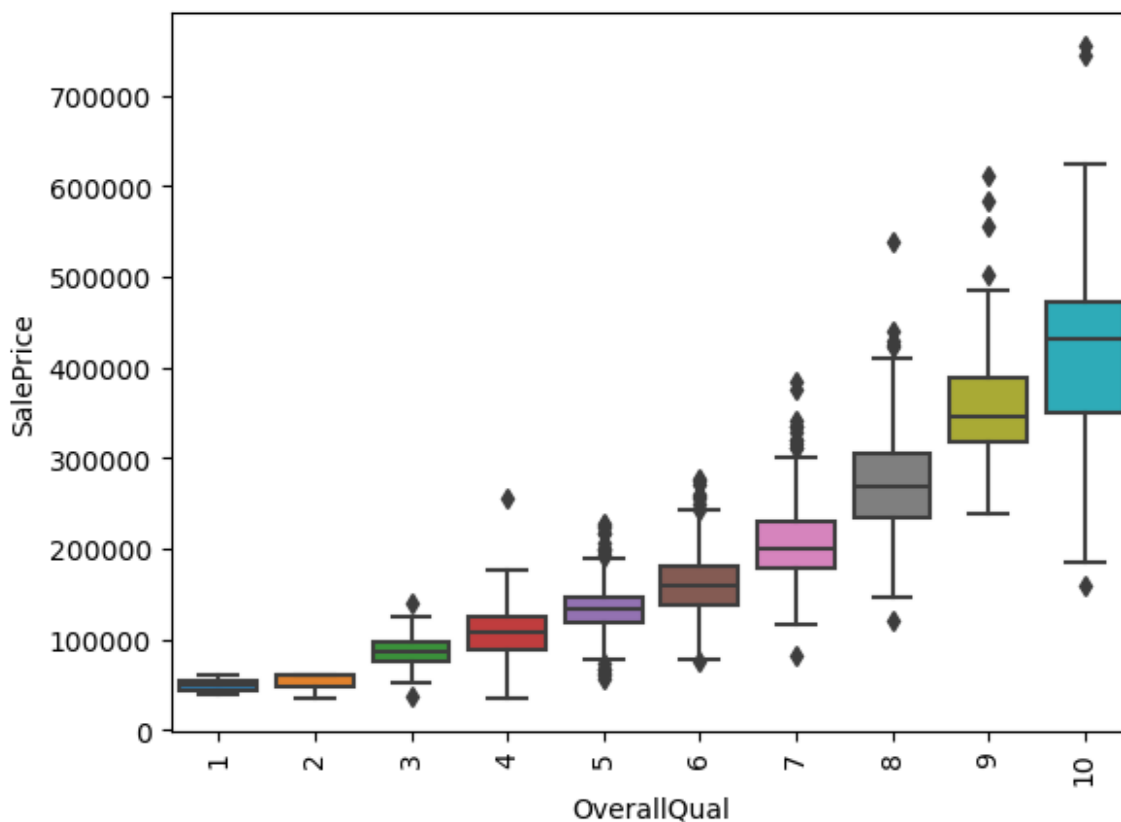
Let's create a bar chart of quantitative correlations.

```
corr_matrix_sorted.plot(kind='bar')  
<Axes: >
```



Predictor 1: 'OverallQual'

```
plot = sns.boxplot(data=train_df, x='OverallQual', y='SalePrice')  
plot.set_xticklabels(plot.get_xticklabels(), rotation=90)  
plt.show()
```

```
train_df.groupby(['OverallQual']).SalePrice.median().sort_values(ascending=False)
```

OverallQual

10 432390.0

9 345000.0

8 269750.0

7 200141.0

6 160000.0

5 133000.0

4 108000.0

3 86250.0

2 60000.0

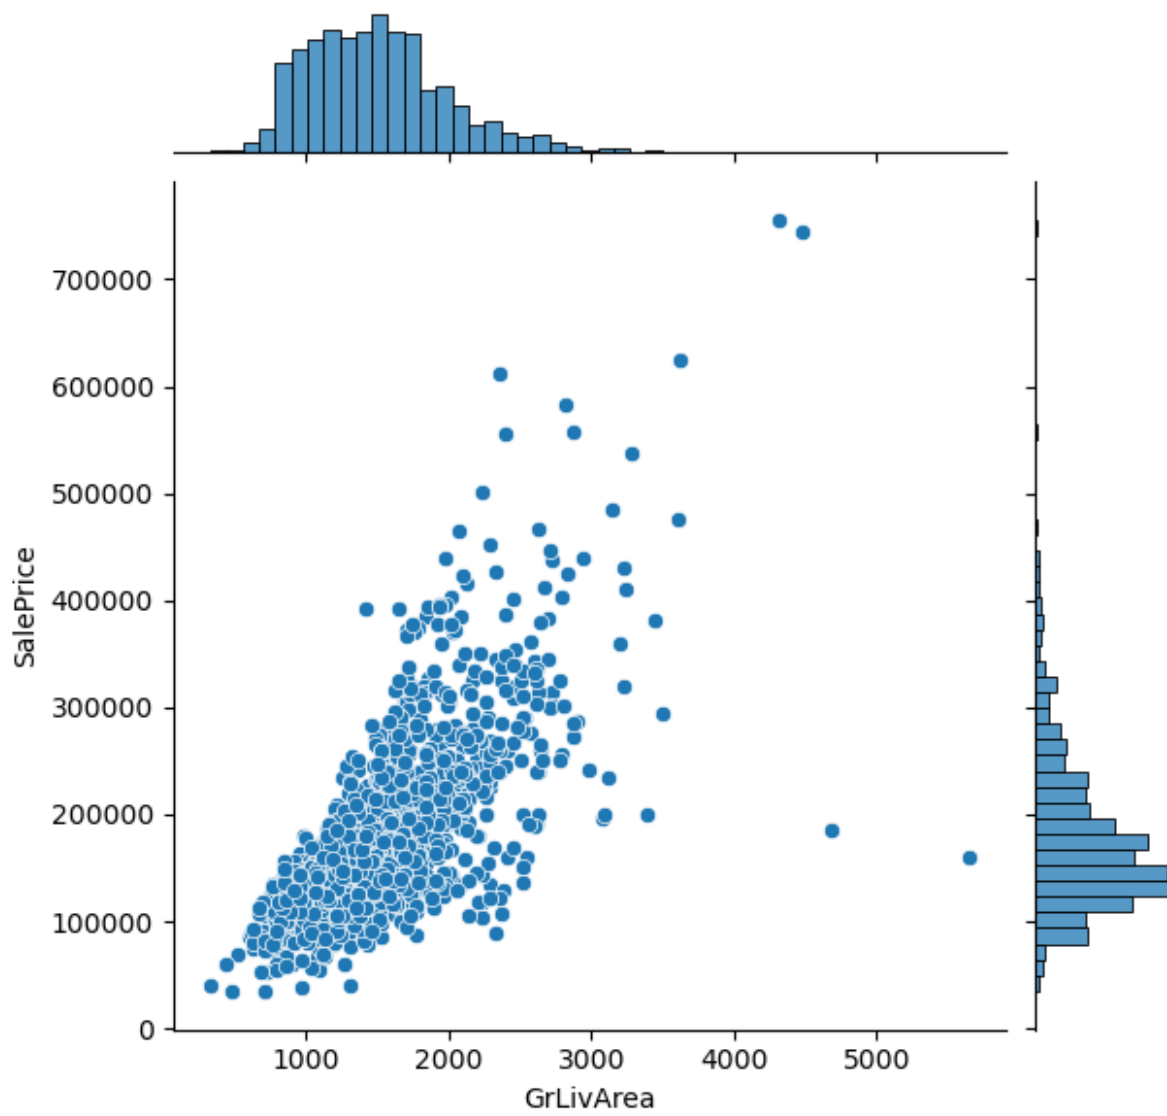
1 50150.0

Name: SalePrice, dtype: float64

The data indicates that as the **OverallQual** rating increases, the median **SalePrice** of homes rises significantly, with top-quality homes (rating 10) having a median price of 432,390 and the lowest quality homes (rating 1) having a median price of 50,150.

Predictor 2: 'GrLivArea'

```
sns.jointplot(x="GrLivArea", y="SalePrice", data=train_df);
```



```
train_df.groupby('GrLivArea').SalePrice.describe()
```

	count	mean	std	min	25%	50%
75% \						
GrLivArea						
334	1.0	39300.0	NaN	39300.0	39300.0	39300.0
39300.0						
438	1.0	60000.0	NaN	60000.0	60000.0	60000.0
60000.0						
480	1.0	35311.0	NaN	35311.0	35311.0	35311.0
35311.0						
520	1.0	68500.0	NaN	68500.0	68500.0	68500.0
68500.0						
605	1.0	86000.0	NaN	86000.0	86000.0	86000.0

```

86000.0
...
.
3627      1.0  625000.0  NaN  625000.0  625000.0  625000.0
625000.0
4316      1.0  755000.0  NaN  755000.0  755000.0  755000.0
755000.0
4476      1.0  745000.0  NaN  745000.0  745000.0  745000.0
745000.0
4676      1.0  184750.0  NaN  184750.0  184750.0  184750.0
184750.0
5642      1.0  160000.0  NaN  160000.0  160000.0  160000.0
160000.0

max
GrLivArea
334      39300.0
438      60000.0
480      35311.0
520      68500.0
605      86000.0
...
3627      625000.0
4316      755000.0
4476      745000.0
4676      184750.0
5642      160000.0

[861 rows x 8 columns]

```

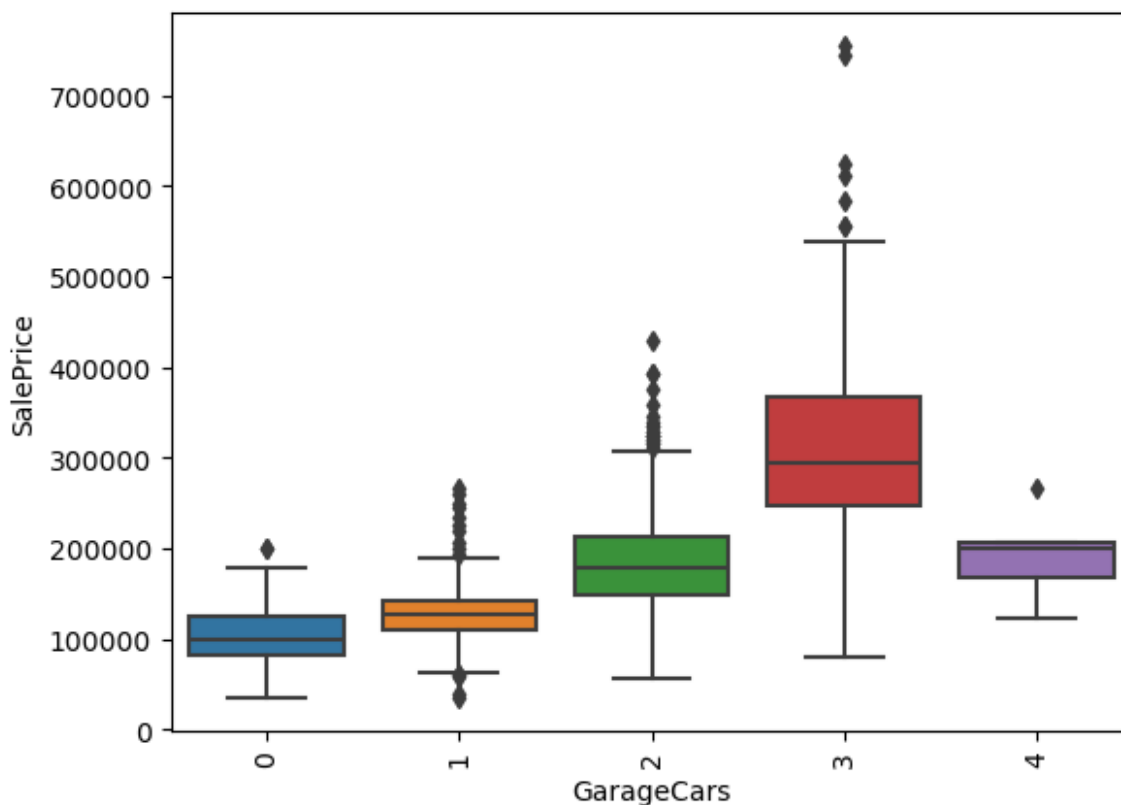
Based on the observation that GrLivingArea (above ground living area) correlates well with SalePrice, with the exception of two larger homes that have a SalePrice just over the median or average, it indicates that while overall there is a strong positive relationship between living area and price, some larger homes may be underpriced or not following the general trend.

Predictor 3: 'GarageCars'

```

plot = sns.boxplot(data=train_df, x='GarageCars', y='SalePrice')
plot.set_xticklabels(plot.get_xticklabels(), rotation = 90)
plt.show()

```



The summary statistics show that as the number of GarageCars increases, the median and mean SalePrice generally rise, with homes having 3 GarageCars having the highest median and mean SalePrice among the categories.

```
train_df.groupby('GarageCars').SalePrice.describe()
```

	count	mean	std	min	25%
50% \					
GarageCars					
0	81.0	103317.283951	32815.023389	34900.0	82500.0
100000.0					
1	369.0	128116.688347	30412.386890	35311.0	110000.0
128000.0					
2	824.0	183851.663835	51617.144258	55993.0	148000.0
177750.0					
3	181.0	309636.121547	106832.925939	81000.0	246578.0
295000.0					
4	5.0	192655.800000	52621.839745	123000.0	168000.0
200000.0					
	75%	max			
GarageCars					
0	124000.0	200500.0			
1	142000.0	266500.0			

2	213000.0	430000.0
3	367294.0	755000.0
4	206300.0	265979.0

The median SalePrice increases as the number of GarageCars increases. Homes with 3 GarageCars have the highest median SalePrice of 295,000, followed by those with 4 GarageCars at 200,000. Homes with 2 GarageCars have a median SalePrice of 177,750, while those with 1 GarageCar have a median SalePrice of 128,000. Homes without a garage (0 GarageCars) have the lowest median SalePrice at 100,000

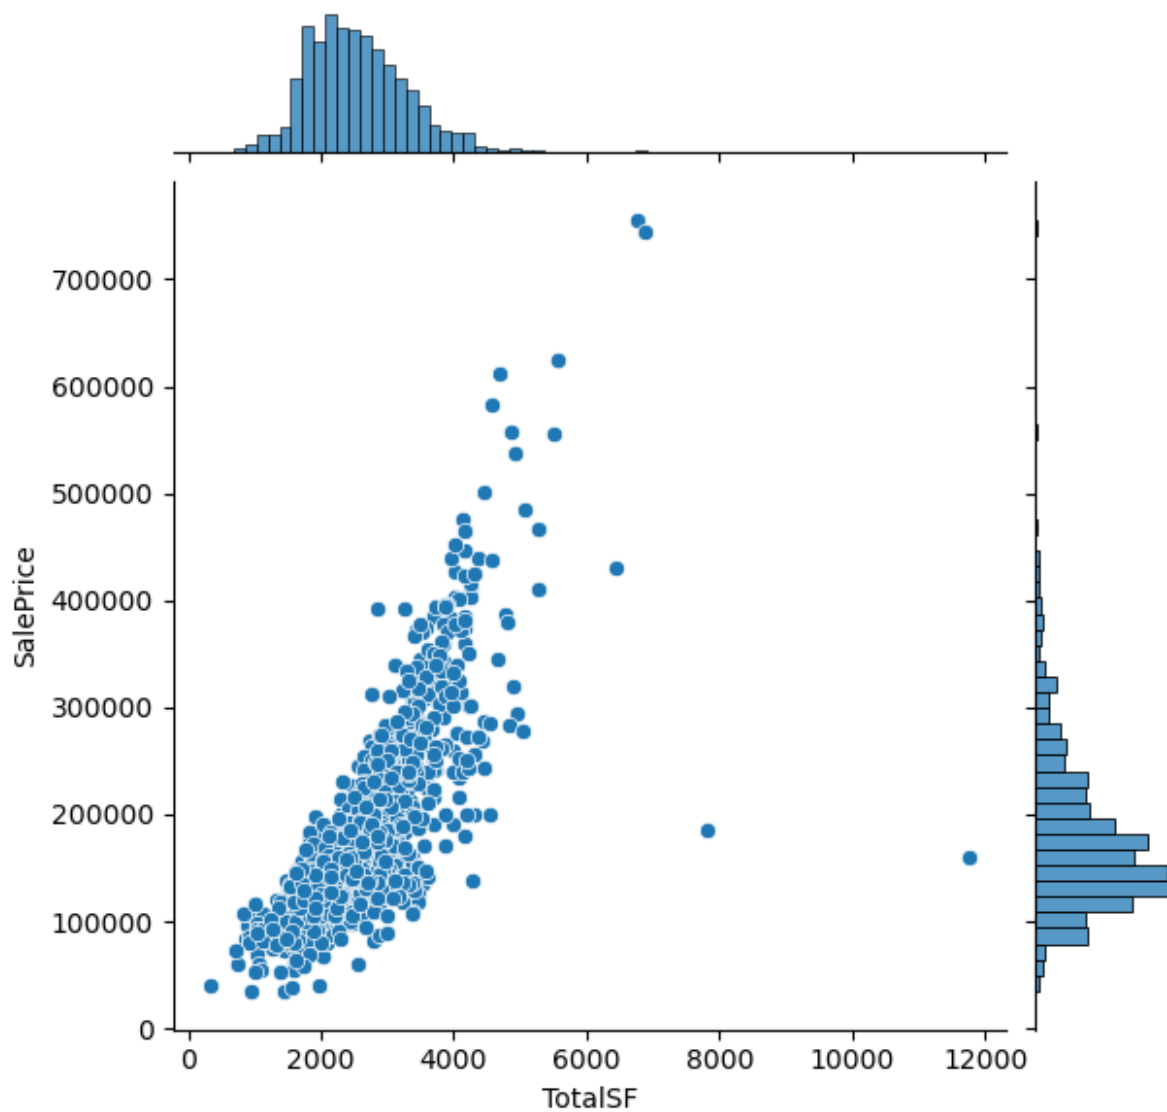
```
train_df.groupby(['GarageCars']).SalePrice.median().sort_values(ascending=False)
```

```
GarageCars
3      295000.0
4      200000.0
2      177750.0
1      128000.0
0      100000.0
Name: SalePrice, dtype: float64
```

4. Engage in feature creation by splitting, merging, or otherwise generating a new predictor.

New features may enable us to create more accurate prediction models for home sale prices. Let's create new attribute 'TotalSF' which is TotalBsmtSF, 1stFlrSF, and 2ndFlrSF combined.

```
train_df['TotalSF'] = train_df['TotalBsmtSF'] + train_df['1stFlrSF'] + train_df['2ndFlrSF']
sns.jointplot(x="TotalSF", y="SalePrice", data=train_df)
<seaborn.axisgrid.JointGrid at 0x152a28cd0>
```



```
train_df_test = train_df[['TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
                           'TotalSF', 'GrLivArea', 'SalePrice']].copy()
train_df_test
```

	TotalBsmtSF	1stFlrSF	2ndFlrSF	TotalSF	GrLivArea	SalePrice
0	856	856	854	2566	1710	208500
1	1262	1262	0	2524	1262	181500
2	920	920	866	2706	1786	223500
3	756	961	756	2473	1717	140000
4	1145	1145	1053	3343	2198	250000
...
1455	953	953	694	2600	1647	175000
1456	1542	2073	0	3615	2073	210000
1457	1152	1188	1152	3492	2340	266500
1458	1078	1078	0	2156	1078	142125

```
1459          1256          1256          0          2512          1256          147500
[1460 rows x 6 columns]
```

5. Using the dependent variable, perform both min-max and standard scaling in Python.

The focus of this section appears to be exclusively on scaling the dependent variable, not the independent variables. **MinMax Scaling: comparing GrLivArea and OverallQual in respect to the dependent variable (SalePrice)**

```
features = ['GrLivArea', 'OverallQual']
features_df = train_df[features]

# Scale the features using MinMaxScaler
scaler = MinMaxScaler()
minmax_features = scaler.fit_transform(features_df)
sale_price = train_df["SalePrice"]

# Create a scatter plot
fig, ax = plt.subplots(figsize=(10, 5))

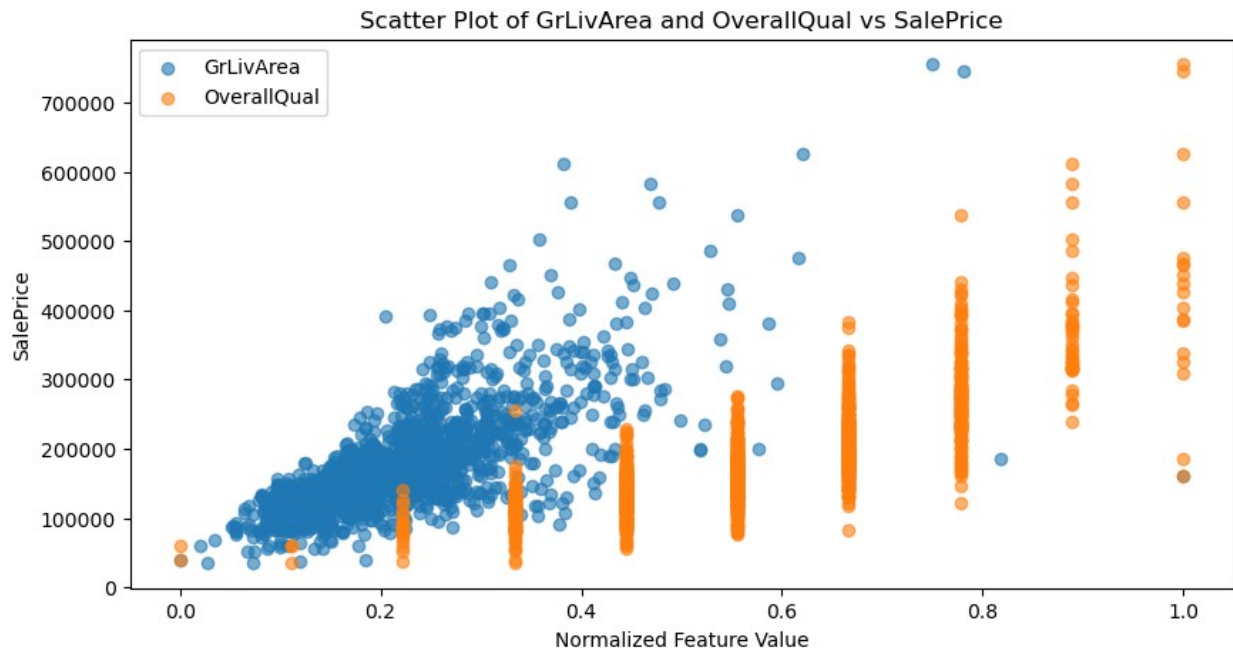
# Scatter plot for GrLivArea
ax.scatter(x=minmax_features[:, 0], y=sale_price, label='GrLivArea',
alpha=0.6)

# Scatter plot for OverallQual
ax.scatter(x=minmax_features[:, 1], y=sale_price, label='OverallQual',
alpha=0.6)

# Add legend
ax.legend()

# Add labels and title
ax.set_xlabel('Normalized Feature Value')
ax.set_ylabel('SalePrice')
ax.set_title('Scatter Plot of GrLivArea and OverallQual vs SalePrice')

# Display the plot
plt.show()
```



MinMax scaling combined with logarithmic transformation helps align the distribution of TotalSF more closely with OverallQual's general distribution, thereby reducing the impact of outliers.

```
features = ['GrLivArea', 'OverallQual']
features_df = train_df[features]

# Log-transform the features
features_df_log = np.log(features_df)

# Scale the log-transformed features using MinMaxScaler
scaler = MinMaxScaler()
minmax_features = scaler.fit_transform(features_df_log)

# Log-transform the SalePrice
sale_price_log = np.log(train_df["SalePrice"])

# Create a scatter plot
fig, ax = plt.subplots(figsize=(10, 5))

# Scatter plot for GrLivArea
ax.scatter(x=minmax_features[:, 0], y=sale_price_log,
          label='GrLivArea', alpha=0.6)

# Scatter plot for OverallQual
ax.scatter(x=minmax_features[:, 1], y=sale_price_log,
          label='OverallQual', alpha=0.6)

# Add legend
ax.legend()
```

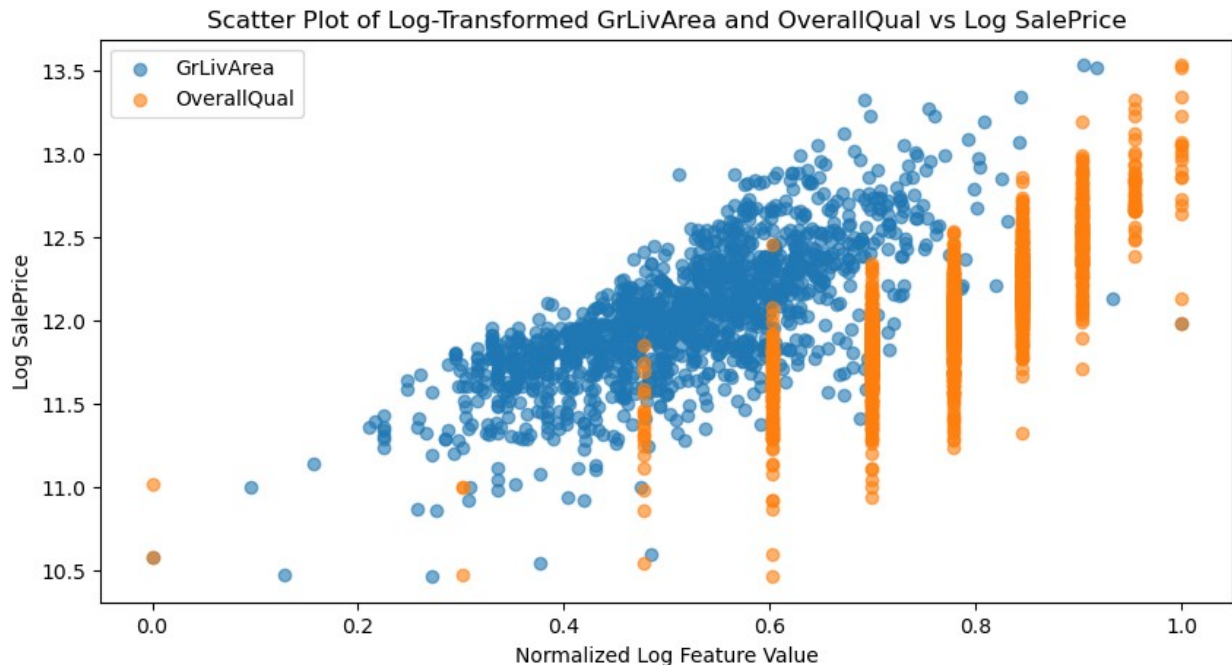


```

# Add labels and title
ax.set_xlabel('Normalized Log Feature Value')
ax.set_ylabel('Log SalePrice')
ax.set_title('Scatter Plot of Log-Transformed GrLivArea and OverallQual vs Log SalePrice')

# Display the plot
plt.show()

```



Standard Scaling: comparing TotalSF and OverallQual in respect to the SalePrice

```

features = ['GrLivArea', 'OverallQual']
features_df = train_df[features]

# Standardize the features using StandardScaler
scaler = StandardScaler()
standardized_features = scaler.fit_transform(features_df)

# SalePrice
sale_price = train_df["SalePrice"]

# Create a scatter plot
fig, ax = plt.subplots(figsize=(10, 5))

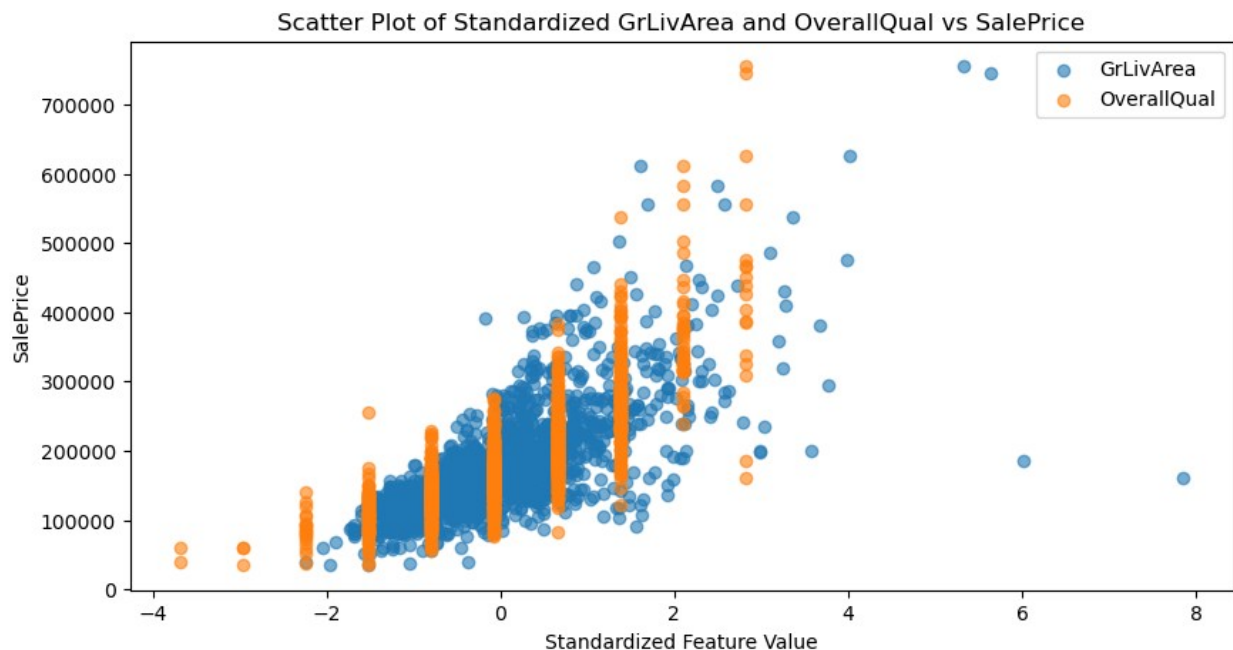
# Plot each feature against SalePrice
for idx, val in enumerate(features):
    ax.scatter(x=standardized_features[:, idx], y=sale_price,
              label=val, alpha=0.6)

```

```
# Add legend
ax.legend()

# Add labels and title
ax.set_xlabel('Standardized Feature Value')
ax.set_ylabel('SalePrice')
ax.set_title('Scatter Plot of Standardized GrLivArea and OverallQual vs SalePrice')

# Display the plot
plt.show()
```



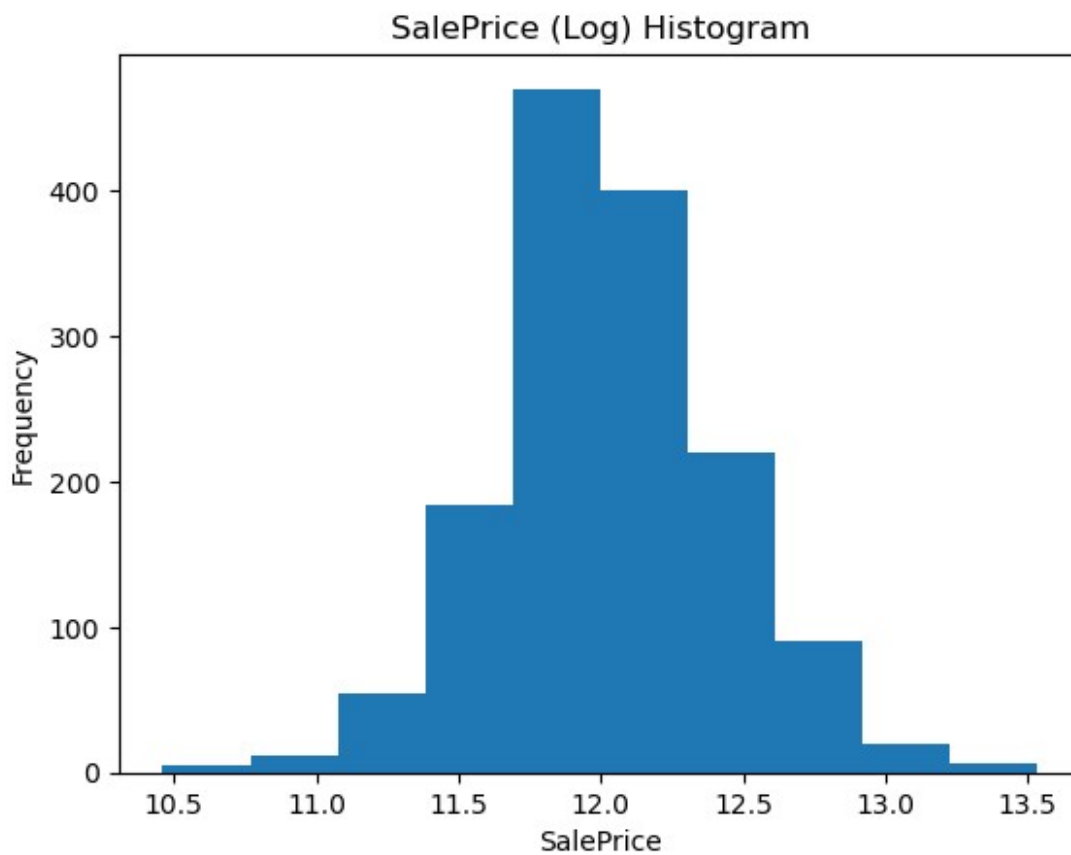
Standard scaling combined with logarithmic transformation doesn't have as dramatic an effect as MinMax scaling with logarithmic transformation because the distributions of the features were already fairly similar before transformation.

```
# General Min max data and general statistics
print("SalePrice Statistics\n")
print("Min house price: ${:,}".format(np.min(train_df['SalePrice'])))
print("Median house price $
{:,}".format(np.median(train_df['SalePrice'])))
print("Max house price: ${:,}".format(np.max(train_df['SalePrice'])))
print("Mean house price: $
{:,}".format(np.mean(train_df['SalePrice'])))
print("Standard deviation of prices: $
{:,}".format(np.std(train_df['SalePrice'])))
```

SalePrice Statistics

Min house price: \$34,900
Median house price \$163,000.0
Max house price: \$755,000
Mean house price: \$180,921.19589041095
Standard deviation of prices: \$79,415.29188606751

```
log_transformed = np.log1p(train_df['SalePrice'])  
plt.hist(log_transformed)  
plt.title('SalePrice (Log) Histogram')  
plt.xlabel('SalePrice')  
plt.ylabel('Frequency')  
plt.show()  
print("Skewness: %f" % log_transformed.skew())
```



Skewness: 0.121347

(N/A)