

# Company Bankruptcy Prediction Report

**Introduction:** Company bankruptcy prediction plays a vital role in financial decision-making, helping stakeholders assess the financial health and stability of businesses. This report explores the effectiveness of SVM, Logistic Regression, and Naive Bayes models in predicting bankruptcy using financial data from the Taiwan Economic Journal.

**Methodology:** We utilized a dataset from 1999 to 2009, encompassing 95 financial indicators, to train and evaluate our models. The data was split into training and validation sets, with preprocessing steps including feature scaling and oversampling to handle imbalance.

## Results:

### Support Vector Machine (SVM):

- **Accuracy:** 88.83%
- **Precision and Recall:** Achieved balanced performance with approximately 91% precision and recall for both bankrupt and non-bankrupt classes.
- **Confusion Matrix:** Showed 1154 true negatives and 1191 true positives, indicating effective discrimination between classes.

### Logistic Regression:

- **Accuracy:** Disappointingly low at 35.61%, indicating poor predictive performance.
- **Precision and Recall:** Particularly weak recall for bankrupt class (8%), highlighting challenges in identifying financially distressed companies.
- **Challenges:** The model struggled with class imbalance and may require further feature engineering or alternative approaches.

### Naive Bayes:

- **Accuracy:** 53.67%
- **Precision and Recall:** Demonstrated higher recall (93%) for bankrupt class but at the expense of precision, indicating a tendency to misclassify non-bankrupt companies.
- **Performance:** Moderate effectiveness in predicting bankruptcies, with an F1-score of 66% for bankrupt class.

**Conclusion:** The SVM model outperformed both Logistic Regression and Naive Bayes in predicting company bankruptcies based on financial indicators. Its balanced precision, recall, and overall accuracy make it a reliable tool for stakeholders looking to assess financial risk. However, the Logistic Regression and Naive Bayes models showed limitations in handling the complexities of the dataset, especially in dealing with class imbalance and predicting bankruptcies accurately.

#### **Recommendations:**

- **Feature Engineering:** Further exploration of feature selection and engineering techniques could enhance model performance, especially for Logistic Regression and Naive Bayes.
- **Ensemble Methods:** Consider ensemble methods such as Random Forest to leverage the strengths of different models and improve predictive accuracy.
- **Continuous Monitoring:** Implement a system for continuous monitoring and updating of predictive models with real-time financial data to adapt to changing business conditions.

**Implications:** Effective bankruptcy prediction models are crucial for investors, creditors, and other stakeholders to make informed decisions and mitigate financial risks. As financial markets evolve, robust predictive analytics will remain essential for ensuring sustainable business practices and economic stability.

# Module 4 Assignment 2 - Company Bankruptcy Prediction (Kaggle)

Sachin Sharma

MSDS-422

07/13/2024

## Management/Research Question

In layman's terms, what is the management/research question of interest, and why would anyone care?

## Requirements

Split the training set into an 80% training and 20% validation set and conduct / improve upon previous EDA. Build at least three models: an SVM, a logistic regression model, a Naïve Bayes model. Evaluate each of the models' assumptions. Conduct hyperparameter tuning for the SVM kernel. Evaluate goodness of fit metrics including TPR, FPR, precision, recall, and accuracy on the training and validation sets. Build ROC and Precision / Recall graphs. Evaluate your models' performance on the validation set using the F1-score. Python scikit-learn should be your primary environment for conducting this research.

## Libraries to be loaded:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold, train_test_split
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report,
precision_recall_curve, roc_curve, accuracy_score
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import RFE
```

## Read Data into Pandas DF

```
train_df = pd.read_csv("data.csv")
train_df.head()
```

	Bankrupt?	ROA(C) before interest and depreciation before interest
\		
0	1	0.370594
1	1	0.464291
2	1	0.426071
3	1	0.399844
4	1	0.465022
	ROA(A) before interest and % after tax	\
0		0.424389
1		0.538214
2		0.499019
3		0.451265
4		0.538432
	ROA(B) before interest and depreciation after tax	\
0		0.405750
1		0.516730
2		0.472295
3		0.457733
4		0.522298
	Operating Gross Margin	Realized Sales Gross Margin \
0	0.601457	0.601457
1	0.610235	0.610235
2	0.601450	0.601364
3	0.583541	0.583541
4	0.598783	0.598783
	Operating Profit Rate	Pre-tax net Interest Rate \
0	0.998969	0.796887
1	0.998946	0.797380
2	0.998857	0.796403
3	0.998700	0.796967
4	0.998973	0.797366
	After-tax net Interest Rate	Non-industry income and expenditure/revenue \
0	0.808809	
0.302646		
1	0.809301	
0.303556		
2	0.808388	
0.302035		
3	0.808966	
0.303350		

4 0.809304  
0.303475

	...	Net Income to Total Assets	Total assets to GNP price \
0	...	0.716845	0.009219
1	...	0.795297	0.008323
2	...	0.774670	0.040003
3	...	0.739555	0.003252
4	...	0.795016	0.003878

	No-credit Interval	Gross Profit to Sales \
0	0.622879	0.601453
1	0.623652	0.610237
2	0.623841	0.601449
3	0.622929	0.583538
4	0.623521	0.598782

	Net Income to Stockholder's Equity	Liability to Equity \
0	0.827890	0.290202
1	0.839969	0.283846
2	0.836774	0.290189
3	0.834697	0.281721
4	0.839973	0.278514

	Degree of Financial Leverage (DFL) \
0	0.026601
1	0.264577
2	0.026555
3	0.026697
4	0.024752

	Interest Coverage Ratio (Interest expense to EBIT)	Net Income Flag \
0	0.564050	
1		
1	0.570175	
1		
2	0.563706	
1		
3	0.564663	
1		
4	0.575617	
1		

	Equity to Liability
0	0.016469
1	0.020794
2	0.016474
3	0.023982
4	0.035490

```
[5 rows x 96 columns]
```

```
train_df.shape
```

```
(6819, 96)
```

```
train_df.describe()
```

	Bankrupt?	ROA(C) before interest and depreciation before interest \
count	6819.000000	6819.000000
mean	0.032263	0.505180
std	0.176710	0.060686
min	0.000000	0.000000
25%	0.000000	0.476527
50%	0.000000	0.502706
75%	0.000000	0.535563
max	1.000000	1.000000

	ROA(A) before interest and % after tax \
count	6819.000000
mean	0.558625
std	0.065620
min	0.000000
25%	0.535543
50%	0.559802
75%	0.589157
max	1.000000

	ROA(B) before interest and depreciation after tax \
count	6819.000000
mean	0.553589
std	0.061595
min	0.000000
25%	0.527277
50%	0.552278
75%	0.584105
max	1.000000

	Operating Gross Margin	Realized Sales Gross Margin \
count	6819.000000	6819.000000
mean	0.607948	0.607929

std	0.016934	0.016916
min	0.000000	0.000000
25%	0.600445	0.600434
50%	0.605997	0.605976
75%	0.613914	0.613842
max	1.000000	1.000000

	Operating Profit Rate	Pre-tax net Interest Rate \
count	6819.000000	6819.000000
mean	0.998755	0.797190
std	0.013010	0.012869
min	0.000000	0.000000
25%	0.998969	0.797386
50%	0.999022	0.797464
75%	0.999095	0.797579
max	1.000000	1.000000

	After-tax net Interest Rate \
count	6819.000000
mean	0.809084
std	0.013601
min	0.000000
25%	0.809312
50%	0.809375
75%	0.809469
max	1.000000

	Non-industry income and expenditure/revenue ... \
count	6819.000000 ...
mean	0.303623 ...
std	0.011163 ...
min	0.000000 ...
25%	0.303466 ...
50%	0.303525 ...
75%	0.303585 ...
max	1.000000 ...

	Net Income to Total Assets	Total assets to GNP price \
count	6819.000000	6.819000e+03
mean	0.807760	1.862942e+07
std	0.040332	3.764501e+08
min	0.000000	0.000000e+00
25%	0.796750	9.036205e-04
50%	0.810619	2.085213e-03
75%	0.826455	5.269777e-03
max	1.000000	9.820000e+09

	No-credit Interval	Gross Profit to Sales \
count	6819.000000	6819.000000
mean	0.623915	0.607946

std	0.012290	0.016934
min	0.000000	0.000000
25%	0.623636	0.600443
50%	0.623879	0.605998
75%	0.624168	0.613913
max	1.000000	1.000000

	Net Income to Stockholder's Equity	Liability to Equity \
count	6819.000000	6819.000000
mean	0.840402	0.280365
std	0.014523	0.014463
min	0.000000	0.000000
25%	0.840115	0.276944
50%	0.841179	0.278778
75%	0.842357	0.281449
max	1.000000	1.000000

	Degree of Financial Leverage (DFL) \
count	6819.000000
mean	0.027541
std	0.015668
min	0.000000
25%	0.026791
50%	0.026808
75%	0.026913
max	1.000000

	Interest Coverage Ratio (Interest expense to EBIT) Net Income Flag \
count	6819.000000
6819.0	
mean	0.565358
1.0	
std	0.013214
0.0	
min	0.000000
1.0	
25%	0.565158
1.0	
50%	0.565252
1.0	
75%	0.565725
1.0	
max	1.000000
1.0	

	Equity to Liability
count	6819.000000
mean	0.047578
std	0.050014



min	0.000000
25%	0.024477
50%	0.033798
75%	0.052838
max	1.000000

[8 rows x 96 columns]

## EDA

Get dataframe information:

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 6819 entries, 0 to 6818
```

```
Data columns (total 96 columns):
```

#	Column	Non-Null Count	Dtype
0	Bankrupt?	6819	int64
1	ROA(C) before interest and depreciation before interest	6819	float64
2	ROA(A) before interest and % after tax	6819	float64
3	ROA(B) before interest and depreciation after tax	6819	float64
4	Operating Gross Margin	6819	float64
5	Realized Sales Gross Margin	6819	float64
6	Operating Profit Rate	6819	float64
7	Pre-tax net Interest Rate	6819	float64
8	After-tax net Interest Rate	6819	float64
9	Non-industry income and expenditure/revenue	6819	float64
10	Continuous interest rate (after tax)	6819	float64
11	Operating Expense Rate	6819	float64
12	Research and development expense rate	6819	float64
13	Cash flow rate	6819	float64

14	Interest-bearing debt interest rate	6819
non-null	float64	
15	Tax rate (A)	6819
non-null	float64	
16	Net Value Per Share (B)	6819
non-null	float64	
17	Net Value Per Share (A)	6819
non-null	float64	
18	Net Value Per Share (C)	6819
non-null	float64	
19	Persistent EPS in the Last Four Seasons	6819
non-null	float64	
20	Cash Flow Per Share	6819
non-null	float64	
21	Revenue Per Share (Yuan ¥)	6819
non-null	float64	
22	Operating Profit Per Share (Yuan ¥)	6819
non-null	float64	
23	Per Share Net profit before tax (Yuan ¥)	6819
non-null	float64	
24	Realized Sales Gross Profit Growth Rate	6819
non-null	float64	
25	Operating Profit Growth Rate	6819
non-null	float64	
26	After-tax Net Profit Growth Rate	6819
non-null	float64	
27	Regular Net Profit Growth Rate	6819
non-null	float64	
28	Continuous Net Profit Growth Rate	6819
non-null	float64	
29	Total Asset Growth Rate	6819
non-null	float64	
30	Net Value Growth Rate	6819
non-null	float64	
31	Total Asset Return Growth Rate Ratio	6819
non-null	float64	
32	Cash Reinvestment %	6819
non-null	float64	
33	Current Ratio	6819
non-null	float64	
34	Quick Ratio	6819
non-null	float64	
35	Interest Expense Ratio	6819
non-null	float64	
36	Total debt/Total net worth	6819
non-null	float64	
37	Debt ratio %	6819
non-null	float64	
38	Net worth/Assets	6819

non-null	float64	
39	Long-term fund suitability ratio (A)	6819
non-null	float64	
40	Borrowing dependency	6819
non-null	float64	
41	Contingent liabilities/Net worth	6819
non-null	float64	
42	Operating profit/Paid-in capital	6819
non-null	float64	
43	Net profit before tax/Paid-in capital	6819
non-null	float64	
44	Inventory and accounts receivable/Net value	6819
non-null	float64	
45	Total Asset Turnover	6819
non-null	float64	
46	Accounts Receivable Turnover	6819
non-null	float64	
47	Average Collection Days	6819
non-null	float64	
48	Inventory Turnover Rate (times)	6819
non-null	float64	
49	Fixed Assets Turnover Frequency	6819
non-null	float64	
50	Net Worth Turnover Rate (times)	6819
non-null	float64	
51	Revenue per person	6819
non-null	float64	
52	Operating profit per person	6819
non-null	float64	
53	Allocation rate per person	6819
non-null	float64	
54	Working Capital to Total Assets	6819
non-null	float64	
55	Quick Assets/Total Assets	6819
non-null	float64	
56	Current Assets/Total Assets	6819
non-null	float64	
57	Cash/Total Assets	6819
non-null	float64	
58	Quick Assets/Current Liability	6819
non-null	float64	
59	Cash/Current Liability	6819
non-null	float64	
60	Current Liability to Assets	6819
non-null	float64	
61	Operating Funds to Liability	6819
non-null	float64	
62	Inventory/Working Capital	6819
non-null	float64	

63	Inventory/Current Liability	6819
non-null	float64	
64	Current Liabilities/Liability	6819
non-null	float64	
65	Working Capital/Equity	6819
non-null	float64	
66	Current Liabilities/Equity	6819
non-null	float64	
67	Long-term Liability to Current Assets	6819
non-null	float64	
68	Retained Earnings to Total Assets	6819
non-null	float64	
69	Total income/Total expense	6819
non-null	float64	
70	Total expense/Assets	6819
non-null	float64	
71	Current Asset Turnover Rate	6819
non-null	float64	
72	Quick Asset Turnover Rate	6819
non-null	float64	
73	Working capital Turnover Rate	6819
non-null	float64	
74	Cash Turnover Rate	6819
non-null	float64	
75	Cash Flow to Sales	6819
non-null	float64	
76	Fixed Assets to Assets	6819
non-null	float64	
77	Current Liability to Liability	6819
non-null	float64	
78	Current Liability to Equity	6819
non-null	float64	
79	Equity to Long-term Liability	6819
non-null	float64	
80	Cash Flow to Total Assets	6819
non-null	float64	
81	Cash Flow to Liability	6819
non-null	float64	
82	CF0 to Assets	6819
non-null	float64	
83	Cash Flow to Equity	6819
non-null	float64	
84	Current Liability to Current Assets	6819
non-null	float64	
85	Liability-Assets Flag	6819
non-null	int64	
86	Net Income to Total Assets	6819
non-null	float64	
87	Total assets to GNP price	6819

non-null	float64	
88	No-credit Interval	6819
non-null	float64	
89	Gross Profit to Sales	6819
non-null	float64	
90	Net Income to Stockholder's Equity	6819
non-null	float64	
91	Liability to Equity	6819
non-null	float64	
92	Degree of Financial Leverage (DFL)	6819
non-null	float64	
93	Interest Coverage Ratio (Interest expense to EBIT)	6819
non-null	float64	
94	Net Income Flag	6819
non-null	int64	
95	Equity to Liability	6819
non-null	float64	

dtypes: float64(93), int64(3)  
memory usage: 5.0 MB

There are no null values.

```
train_df.isna().sum().sum()
0

# Separate numeric (quantitative) and categorical (nominal) variables
numeric_vars = train_df.select_dtypes(include=['float64',
'int64']).columns.tolist()
categorical_vars =
train_df.select_dtypes(include=['object']).columns.tolist()

print("\nNumeric Variables:")
print(len(numeric_vars))







































print("\nCategorical Variables:")
print(len(categorical_vars))

Numeric Variables:
96

Categorical Variables:
0
```

Let's plot a histogram of the variables.

```
sns.set(style='whitegrid', font_scale=1.1, rc={'figure.figsize': [30,
102]})
train_df[train_df.columns].hist(bins=15, layout=(24, 4));
```



Let's create a correlation heatmap.

```
correlation = train_df.corr()
fig, ax = plt.subplots(figsize = (15,15));
sns.heatmap(correlation, ax = ax, linewidth = 0.1);
```




Just over three percent of the data contain rows listed as 'Bankrupt?' (the dependent variable) = 1.

Can we mitigate the overwhelming 'Bankrupt?' = 0 bias using some oversampling methodology...?

```
total_bankrupt = train_df['Bankrupt?'].sum()
pct_bankrupt = total_bankrupt/len(train_df['Bankrupt?'])*100
print('Num bankrupt: %d, %% of sample: %.2f%%' %(total_bankrupt,
pct_bankrupt))
```

Num bankrupt: 220, % of sample: 3.23%

```
# Class distribution of the target variable
sns.set(style='whitegrid', font_scale=1.1, rc={"figure.figsize": [8,
5]})
sns.countplot(x='Bankrupt?', data=train_df)
plt.title('Class Distribution of Bankrupt?')
plt.show()
```



```
print('Financially stable:',
round(train_df['Bankrupt?'].value_counts()[0] / len(train_df) * 100,2)
,'%')
print('Financially unstable:',
round(train_df['Bankrupt?'].value_counts()[1] / len(train_df) * 100,
2), '%')
```

Financially stable: 96.77 %  
Financially unstable: 3.23 %



We see the data is highly skewed towards, Financially stable. If we train the model on this dataset, our prediction will be biased towards Financially stable.

We will balance the dataset, to train our model.

```
train_df_X = train_df.copy()
train_df_y = train_df_X['Bankrupt?']
train_df_X.drop(['Bankrupt?'], axis=1, inplace=True)

train_df_X.shape

(6819, 95)
```

## Oversampling

We have the data between 'Bankrupt?' and not 'Bankrupt?' now even using SMOTE oversampling.


```
oversample = SMOTE(random_state=42)
train_df_X, train_df_y = oversample.fit_resample(train_df_X, train_df_y)

total_bankrupt = train_df_y.sum()
pct_bankrupt = total_bankrupt/len(train_df_y)*100
print('Num bankrupt: %d, %% of sample: %.2f%%' %(total_bankrupt,
pct_bankrupt))

Num bankrupt: 6599, % of sample: 50.00%

# Display the new class distribution
print("Resampled class distribution:\n", train_df_y.value_counts())
sns.countplot(x=train_df_y)
plt.title('Resampled Class Distribution')
plt.show()

Resampled class distribution:
Bankrupt?
1      6599
0      6599
Name: count, dtype: int64
```



## Feature Selection

Let's attempt to reduce the 95 independent variables into a smaller subset of more correlated features using a Feature Selection methodology.

We will use Recursive Feature Elimination with the Random Forest Classifier to narrow down from 95 to 15 features.

```
X_train,X_val,y_train,y_val =  
train_test_split(train_df_X,train_df_y,test_size=0.2,random_state=42)  
  
%%time  
select = RFE(RandomForestClassifier(n_estimators=100,  
random_state=42), n_features_to_select=15)  
  
select.fit(X_train, y_train)  
  
mask = select.get_support()  
  
X_train_rfe = select.transform(X_train)  
X_test_rfe = select.transform(X_val)  
  
score = RandomForestClassifier().fit(X_train_rfe,  
y_train).score(X_test_rfe, y_val)
```

```

print("Test score: {:.3f}".format(score), " number of features:
{}".format(15))

features = pd.DataFrame({'features':list(train_df.iloc[:,1:].keys()),
'select':list(mask)})
features = list(features[features['select']==True]['features'])
features.append('Bankrupt?')

Test score: 0.963  number of features: 15
CPU times: user 5min 20s, sys: 309 ms, total: 5min 20s
Wall time: 5min 21s

features = features[:-1]

features

[' Pre-tax net Interest Rate',
' After-tax net Interest Rate',
' Continuous interest rate (after tax)',
' Interest-bearing debt interest rate',
' Persistent EPS in the Last Four Seasons',
' Quick Ratio',
' Interest Expense Ratio',
' Total debt/Total net worth',
' Debt ratio %',
' Borrowing dependency',
' Net profit before tax/Paid-in capital',
' Retained Earnings to Total Assets',
' Cash Turnover Rate',
' Net Income to Total Assets',
' Net Income to Stockholder's Equity"]

train_df_X_reduced = train_df_X[features]
train_df_X_reduced.shape

(13198, 15)

```

## Scale Data

```

scaler = StandardScaler()
train_df_X_reduced_scaled = scaler.fit_transform(train_df_X_reduced)
train_df_X_reduced_scaled =
pd.DataFrame(scaler.transform(train_df_X_reduced),
index=train_df_X_reduced.index, columns=train_df_X_reduced.columns)

train_df_X_reduced_scaled.head(5)

```

	Pre-tax net Interest Rate	After-tax net Interest Rate	\
0	-0.007885	-0.004726	
1	0.044809	0.044772	
2	-0.059589	-0.047220	

3	0.000646	0.011013
4	0.043301	0.045073
Continuous interest rate (after tax) Interest-bearing debt		
interest rate \		
0	-0.018885	-
0.123653		
1	0.037666	-
0.123653		
2	-0.094945	-
0.123653		
3	0.008911	-
0.123653		
4	0.042448	-
0.123653		
Persistent EPS in the Last Four Seasons Quick Ratio \		
0	-1.107635	-0.046312
1	-0.005620	-0.046312
2	-0.790904	-0.046312
3	-0.427056	-0.046312
4	0.093849	-0.046312
Interest Expense Ratio Total debt/Total net worth Debt ratio % \		
0	-0.065798	-0.044923 0.942596
1	0.359610	-0.044923 0.358094
2	-0.091859	-0.044923 0.941625
3	-0.043224	-0.044923 0.041570
4	0.431516	-0.044923 -0.680319
Borrowing dependency Net profit before tax/Paid-in capital \		
0	0.179148	-0.866768
1	-0.140027	0.101639
2	-0.084970	-0.547788
3	-0.069635	-0.562532
4	-0.180963	0.055067
Retained Earnings to Total Assets Cash Turnover Rate \		
0	-0.555638	-0.701870
1	0.346746	0.091643
2	-0.339167	-0.583546
3	-0.436444	-0.087991
4	-0.211240	-0.558944
Net Income to Total Assets Net Income to Stockholder's Equity		
0	-0.896223	-0.117565
1	0.320238	0.150567
2	0.000392	0.079645
3	-0.544098	0.033535
4	0.315877	0.150644

- Split the training set into an 80% training and 20% validation set.

```
# un-scaled split data
X_train,X_val,y_train,y_val=train_test_split(train_df_X_reduced,train_
df_y,test_size=0.2,random_state=42)

# scaled split data
X_train_scaled,X_val_scaled,y_train_scaled,y_val_scaled =
train_test_split(train_df_X_reduced_scaled,train_df_y,test_size=0.2,ra
ndom_state=42)
```

## MODELING

```
# set up a KFold cross-validation rule
K = 10
kf = KFold(n_splits=K, shuffle=True, random_state=42)
```

## Support Vector Machine

```
#training model
svc = SVC(kernel='linear', gamma=0.01, C=2, probability=True)
svc.fit(X_train_scaled,y_train_scaled)

#getting confusion matrix
y_pred = svc.predict(X_val_scaled)
cm = confusion_matrix(y_val_scaled,y_pred)
print('confusion matrix:\n',cm)
print('accuracy score = ',accuracy_score(y_val_scaled,y_pred))
print("Classification Report:\n",classification_report(y_val_scaled,y_pred))
```

```
confusion matrix:
[[1154  180]
 [ 115 1191]]
accuracy score =  0.8882575757575758
Classification Report:
```

	precision	recall	f1-score	support
0	0.91	0.87	0.89	1334
1	0.87	0.91	0.89	1306
accuracy			0.89	2640
macro avg	0.89	0.89	0.89	2640
weighted avg	0.89	0.89	0.89	2640

```
y_pred = svc.predict_proba(X_val_scaled)[::,1]
fpr, tpr, _ = roc_curve(y_val_scaled, y_pred)


#create ROC curve
plt.plot(fpr,tpr, color="blue")
plt.title('ROC Curve - SVM')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.show()


#calculate precision and recall
precision, recall, thresholds = precision_recall_curve(y_val_scaled,
y_pred)

#create precision recall curve
fig, ax = plt.subplots()
ax.plot(recall, precision, color='blue')

#add axis labels to plot
ax.set_title('Precision/Recall Curve - SVM')
ax.set_ylabel('Precision')
ax.set_xlabel('Recall')

#display plot
plt.show()
```





The SVM model applied to the dataset achieved a commendable accuracy of 88.83%. With precision and recall scores around 89% for both bankrupt and non-bankrupt classes, the model demonstrates robust performance in distinguishing between financially stable and distressed companies. The balanced performance metrics indicate that the SVM model effectively utilizes the dataset's features to predict company bankruptcy with high accuracy and reliability.

## Logistic Regression

```
#training model
lr = LogisticRegression(max_iter = 10000)
lr.fit(X_train_scaled,y_train_scaled)

#getting confusion matrix
y_pred = lr.predict(X_val)
cm = confusion_matrix(y_val,y_pred)
print('confusion matrix:\n',cm)
lra = accuracy_score(y_val,y_pred)
print('accuracy score = ',lra)
print("Classification Report:\n",classification_report(y_val,y_pred))

confusion matrix:
[[ 840  494]
 [1206  100]]
```



```
accuracy score = 0.3560606060606061
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.41	0.63	0.50	1334
1	0.17	0.08	0.11	1306
accuracy			0.36	2640
macro avg	0.29	0.35	0.30	2640
weighted avg	0.29	0.36	0.30	2640

```
#define metrics
```

```
y_pred = lr.predict_proba(X_val)[:,:1]
```

```
fpr, tpr, _ = roc_curve(y_val, y_pred)
```

```
#create ROC curve
```


```
plt.plot(fpr,tpr, color="blue")
```

```
plt.title('ROC Curve - Logistic Regression')
```

```
plt.xlabel('False Positive Rate (FPR)')
```

```
plt.ylabel('True Positive Rate (TPR)')
```

```
plt.show()
```



## Precision/Recall Curve

```
# ROC Curve
y_pred_proba = lr.predict_proba(X_val_scaled)[::, 1]
fpr, tpr, _ = roc_curve(y_val, y_pred_proba)


plt.figure()
plt.plot(fpr, tpr, color="blue")
plt.title('ROC Curve - Logistic Regression')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.show()


#calculate precision and recall
precision, recall, thresholds = precision_recall_curve(y_val, y_pred)

#create precision recall curve
fig, ax = plt.subplots()
ax.plot(recall, precision, color='blue')

#add axis labels to plot
ax.set_title('Precision/Recall Curve - Logistic Regression')
ax.set_ylabel('Precision')
ax.set_xlabel('Recall')

#display plot
plt.show()
```





The Logistic Regression model applied to the dataset yielded disappointing results with an accuracy of only 35.61%. The model struggled particularly with recall, achieving 8% for class 1 (bankrupt), indicating a high rate of false negatives. Precision was also low, with values of 41% for class 0 (non-bankrupt) and 17% for class 1, highlighting challenges in correctly identifying bankrupt companies. Overall, the model's performance suggests limitations in effectively leveraging the dataset's features for bankruptcy prediction compared to the SVM model.

## Naive Bayes

```
#training model
nb = GaussianNB()
nb.fit(X_train,y_train)

#getting confusion matrix
y_pred = nb.predict(X_val)
cm = confusion_matrix(y_val,y_pred)
print('confusion matrix:\n',cm)

#checking accuracy
nba = accuracy_score(y_val,y_pred)
print('accuracy score = ',accuracy_score(y_val,y_pred))
print("Classification Report:\n",classification_report(y_val,y_pred))
```

confusion matrix:

```
[[ 207 1127]
 [  96 1210]]
```

accuracy score = 0.5367424242424242

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.16	0.25	1334
1	0.52	0.93	0.66	1306
accuracy			0.54	2640
macro avg	0.60	0.54	0.46	2640
weighted avg	0.60	0.54	0.46	2640

*# ROC Curve*


```
y_pred_proba = nb.predict_proba(X_val)[: , 1]
fpr, tpr, _ = roc_curve(y_val, y_pred_proba)
```


```
plt.figure()
plt.plot(fpr, tpr, color="blue")
plt.title('ROC Curve - Naive Bayes')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.show()
```

*# Precision-Recall Curve (optional)*

```
precision, recall, _ = precision_recall_curve(y_val, y_pred_proba)
```

```
plt.figure()
plt.plot(recall, precision, color='blue')
plt.title('Precision-Recall Curve - Naive Bayes')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.show()
```





The Naive Bayes model applied to the dataset achieved an accuracy of 53.67%. It showed a significant disparity in performance between classes, with a higher recall of 93% for class 1 (bankrupt) compared to 16% for class 0 (non-bankrupt). Precision was also higher for class 1 at 52%, indicating the model's ability to better identify bankrupt companies but at the cost of misclassifying non-bankrupt companies. The overall F1-score was 25% for class 0 and 66% for class 1, reflecting the model's moderate performance in predicting bankruptcies based on the dataset's features.

## CONCLUSION

### Management/Research Question

**In layman's terms, what is the management/research question of interest, and why would anyone care?**

The management/research question of interest could be:

**Research Question:** Can we accurately predict whether a company is likely to go bankrupt based on financial indicators?

**Layman's Explanation:** The question aims to determine if we can use financial information to forecast whether a company might face financial distress and potentially go out of business.

**Why it Matters:** Understanding and predicting company bankruptcy is crucial for various stakeholders, including investors, creditors, and even employees. It helps them make informed

decisions about investments, loans, and employment stability. By identifying early warning signs of financial distress, stakeholders can take proactive measures to mitigate risks or capitalize on opportunities effectively.

